Scaling Language Models

Antoine Bosselut

EPFL

Announcements

- **Assignment 3:** Due Sunday, 30/04/2023 at 11:59 PM
  - Office Hours: **tomorrow**, Thursday 1 PM
- **No Lecture Tomorrow!**
- **Course Project:** Kickoff!
  - Data Packages were released

Next few days: Project Sign-ups

- **To-Dos:**
  - **URGENT:** Look over the Project Description
  - **URGENT:** Fill out team registration form if you haven't already
  - **URGENT:** Get API key to access GPTWrapper Server:
    - Fill out data consent form
    - After filling it out, ML4ED will send API keys
  - **URGENT:** Sign up for project repository
  - Look through README in project repository for details on milestone submission
  - Get started early! **Milestone 1 due May 16th!**

Nostoi

- **Last class,** learning about ChatGPT:
  - Helpful comment that we were connecting a lot of concepts from many past lectures: Transformers (Week 4), Fine-tuning (Week 5), Data quality (Week 6), Text generation (Week 7)
  - Can be difficult to keep track of these many concepts

- **Nostoi: AI-powered slide reader**
  - Cross course search — reference content to older courses
  - Study tools — automatically build flashcards for future studying
  - Highlight words and get personalized explanations (powered by ChatGPT) — **Trust, but verify!**

Today's Outline

- **Lecture**
  - **Quick Recap:** Scale
  - **Managing scale when training:** Scaling laws
- **Guest Lecture:** Reza Banaei
  - **Managing scale when deploying:** Model Compression
    - how can we make LLMs more efficient?

Language Model Scaling

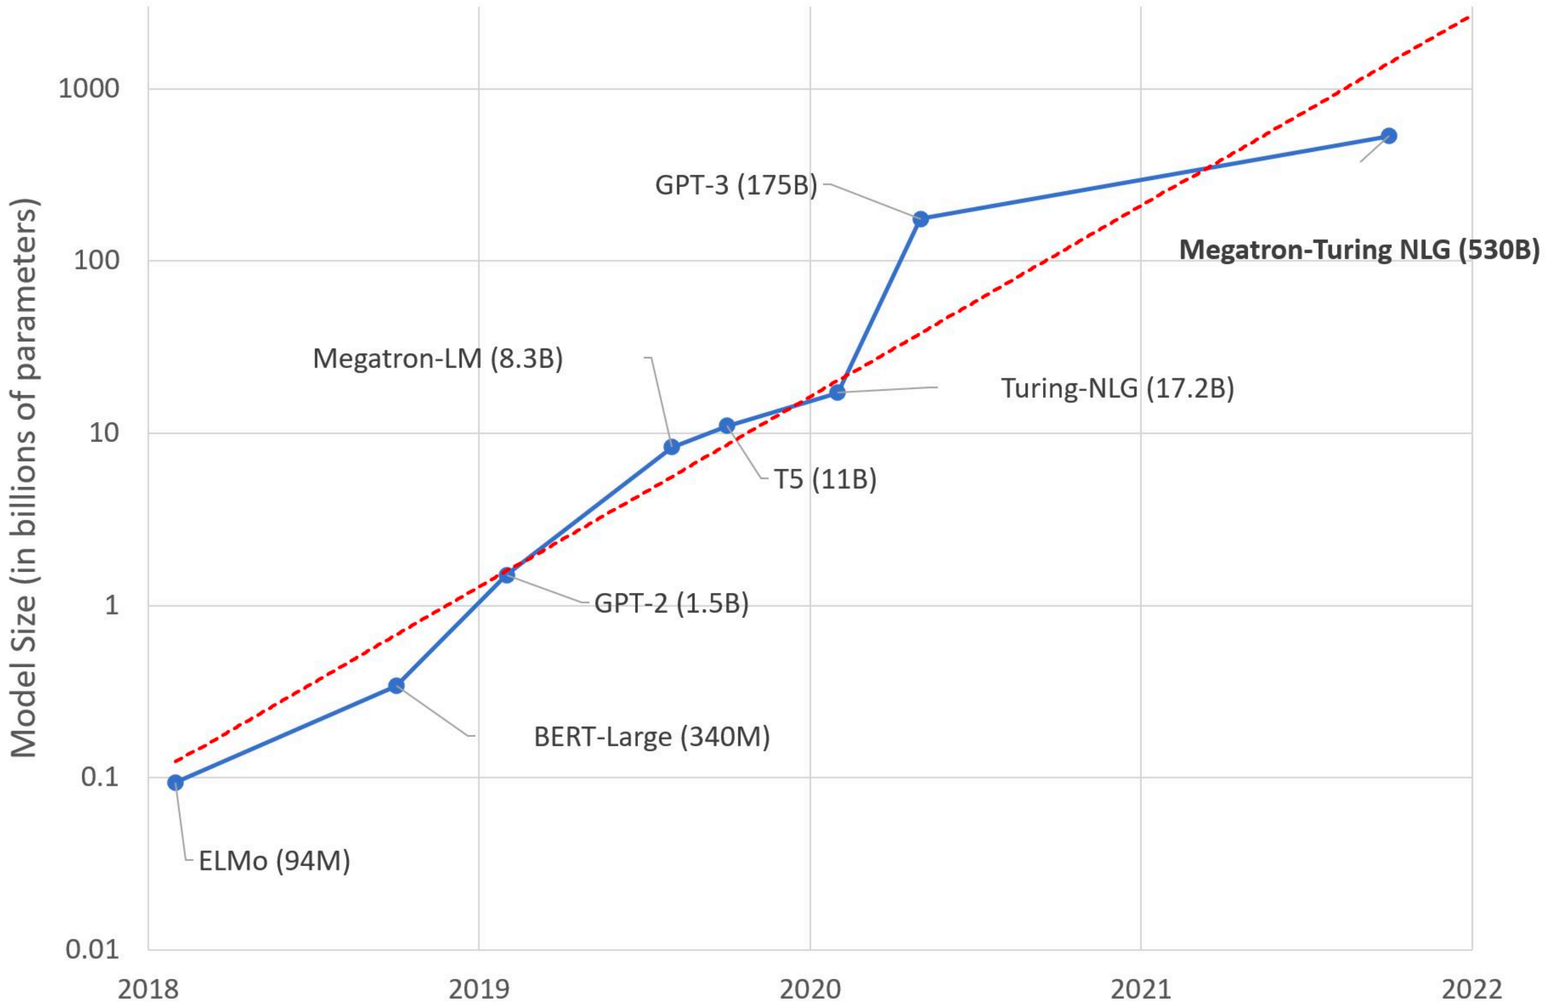

Larger models

More data

More compute

More 

Every part of the model scales!

Model Name	n_{params}	n_{layers}	d_{model}	n_{heads}	d_{head}	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	6.0×10^{-4}
GPT-3 Medium	350M	24	1024	16	64	0.5M	3.0×10^{-4}
GPT-3 Large	760M	24	1536	16	96	0.5M	2.5×10^{-4}
GPT-3 XL	1.3B	24	2048	24	128	1M	2.0×10^{-4}
GPT-3 2.7B	2.7B	32	2560	32	80	1M	1.6×10^{-4}
GPT-3 6.7B	6.7B	32	4096	32	128	2M	1.2×10^{-4}
GPT-3 13B	13.0B	40	5140	40	128	2M	1.0×10^{-4}
GPT-3 175B or “GPT-3”	175.0B	96	12288	96	128	3.2M	0.6×10^{-4}

- Trained on 570 GB of Common Crawl data
- **How?** Used cluster provided by Microsoft

Why scale?

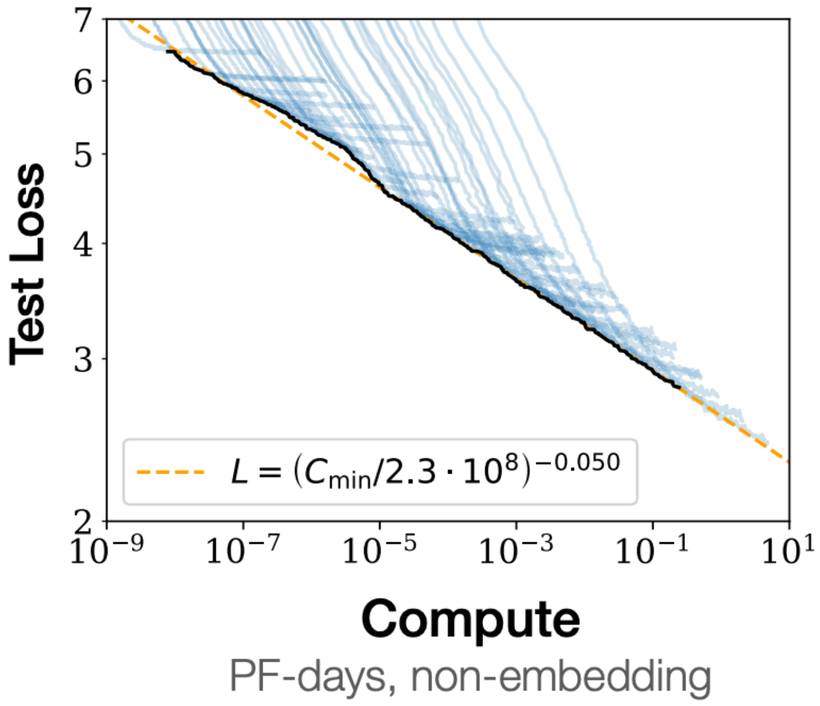

- Last week, we talked about benefits of scaling in terms of **emergence**
- Practically, training for longer also leads to lower test loss
- Larger models can reach lower test losses

Kaplan et al. (2020)

What should we scale?

Model size, dataset size, compute budget

Given a compute budget, how big of a model can we train?
and how big of a dataset should we train it on?

Impact of compute budget

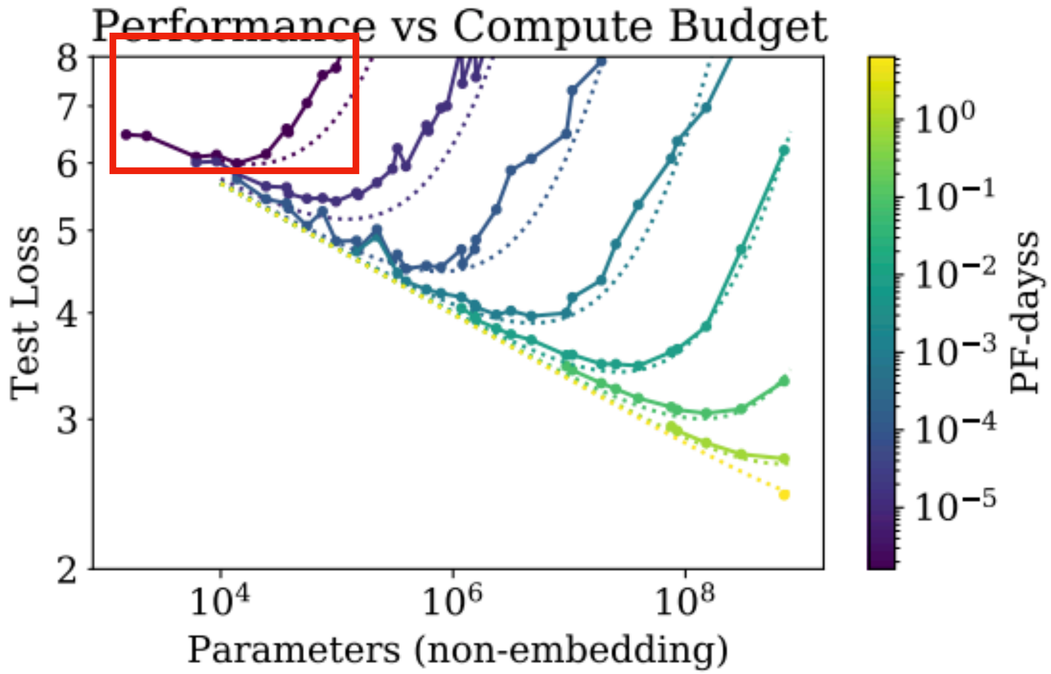

Dotted lines estimate these curves.
Need to predict for larger models!

- For a fixed compute budget, there is an optimal number of parameters that we can train
- Having **too large** a model for **too small** a compute budget does make the model learn
  - Model doesn't see enough examples
- Having **too small** a model for too large a compute budget is also bad
  - Repeated computation isn't helpful if the model has no capacity to encode additional information

Kaplan et al. (2020)

Consideration

- With a fixed compute budget (in FLOP-days), we have two costs:
  - Number of floating point operations needed to train on a single example (model size)
  - Number of total examples we will train on
- **How should we trade off these two costs?**
  - Which should we get prioritise?

Model-Data Trade-off

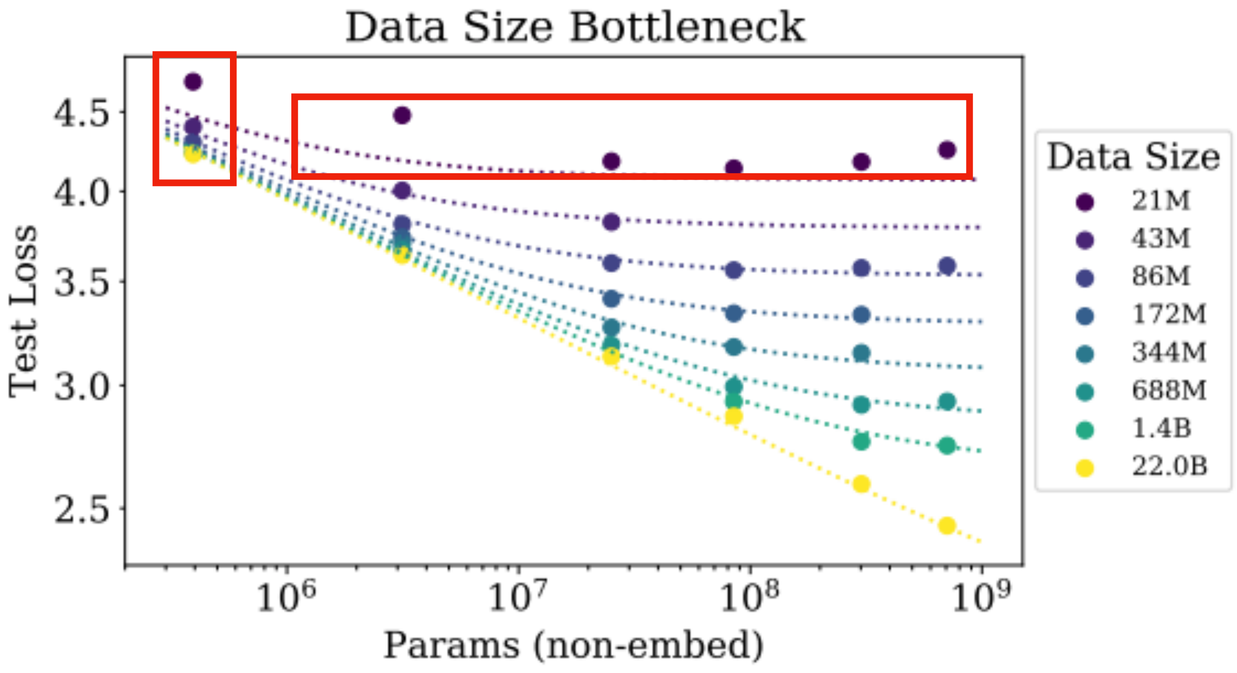

- Larger models benefit more from larger datasets
- **Smaller models saturate**
  - Only so much capacity to learn!
- At some point, larger models don't benefit more from same-sized data
- Data size and model size need to be scaled jointly

Kaplan et al. (2020)

Other Cool Findings

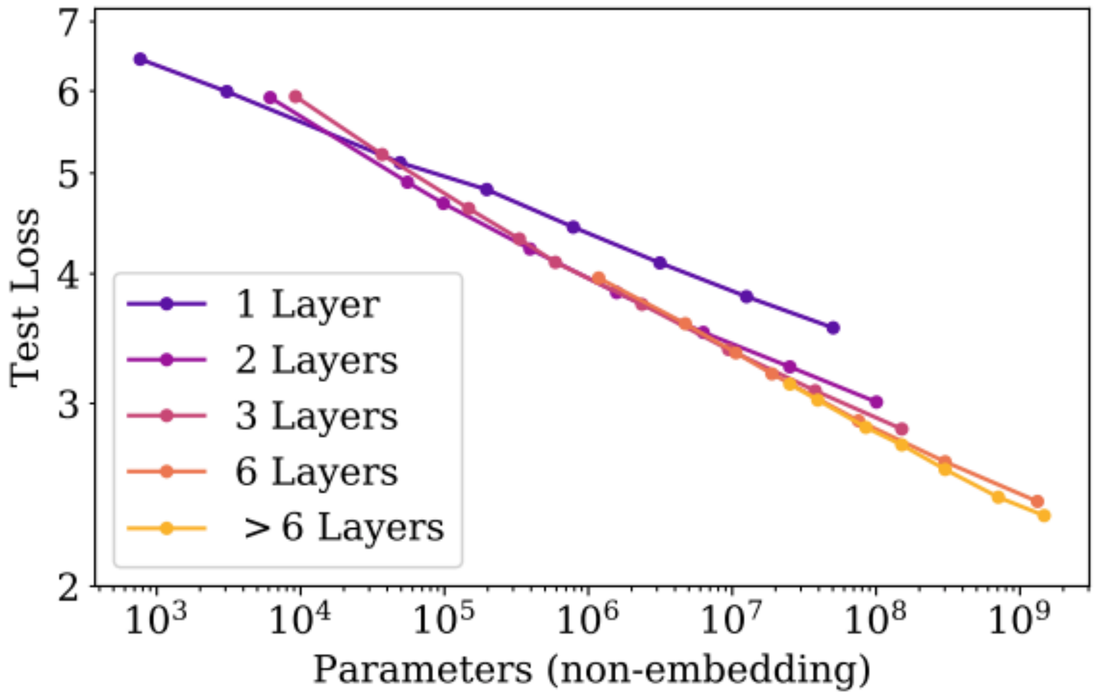

No need to make models terribly deep

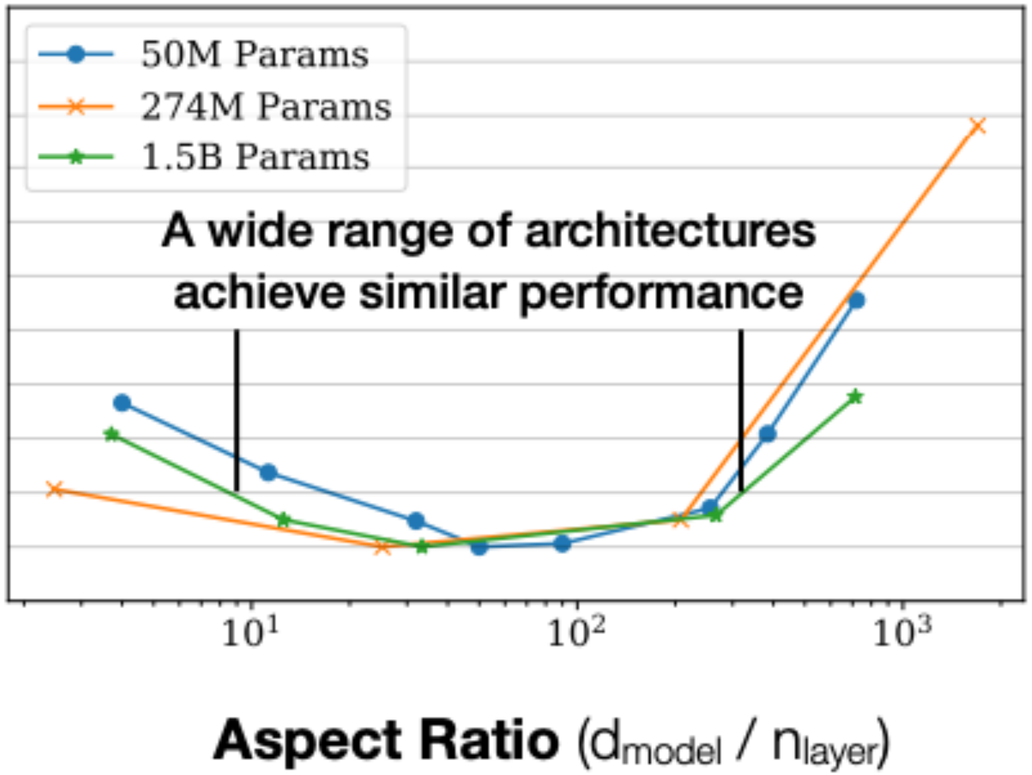

Multiple ratios of depth vs. width (aka embedding size) are fine

Other cool findings

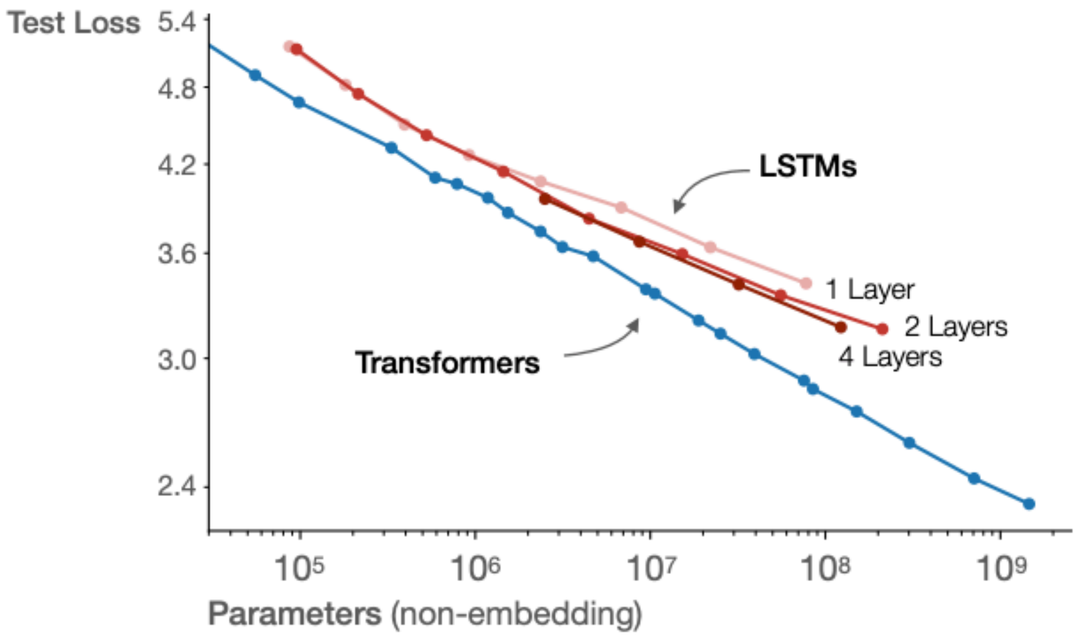

- LSTMs also follow scaling laws, benefitting from increased scale
- They scale less efficiently than transformers, though
- They still have trouble modelling long-term depencies (>100 tokens)

To scale up: estimate model, data, compute

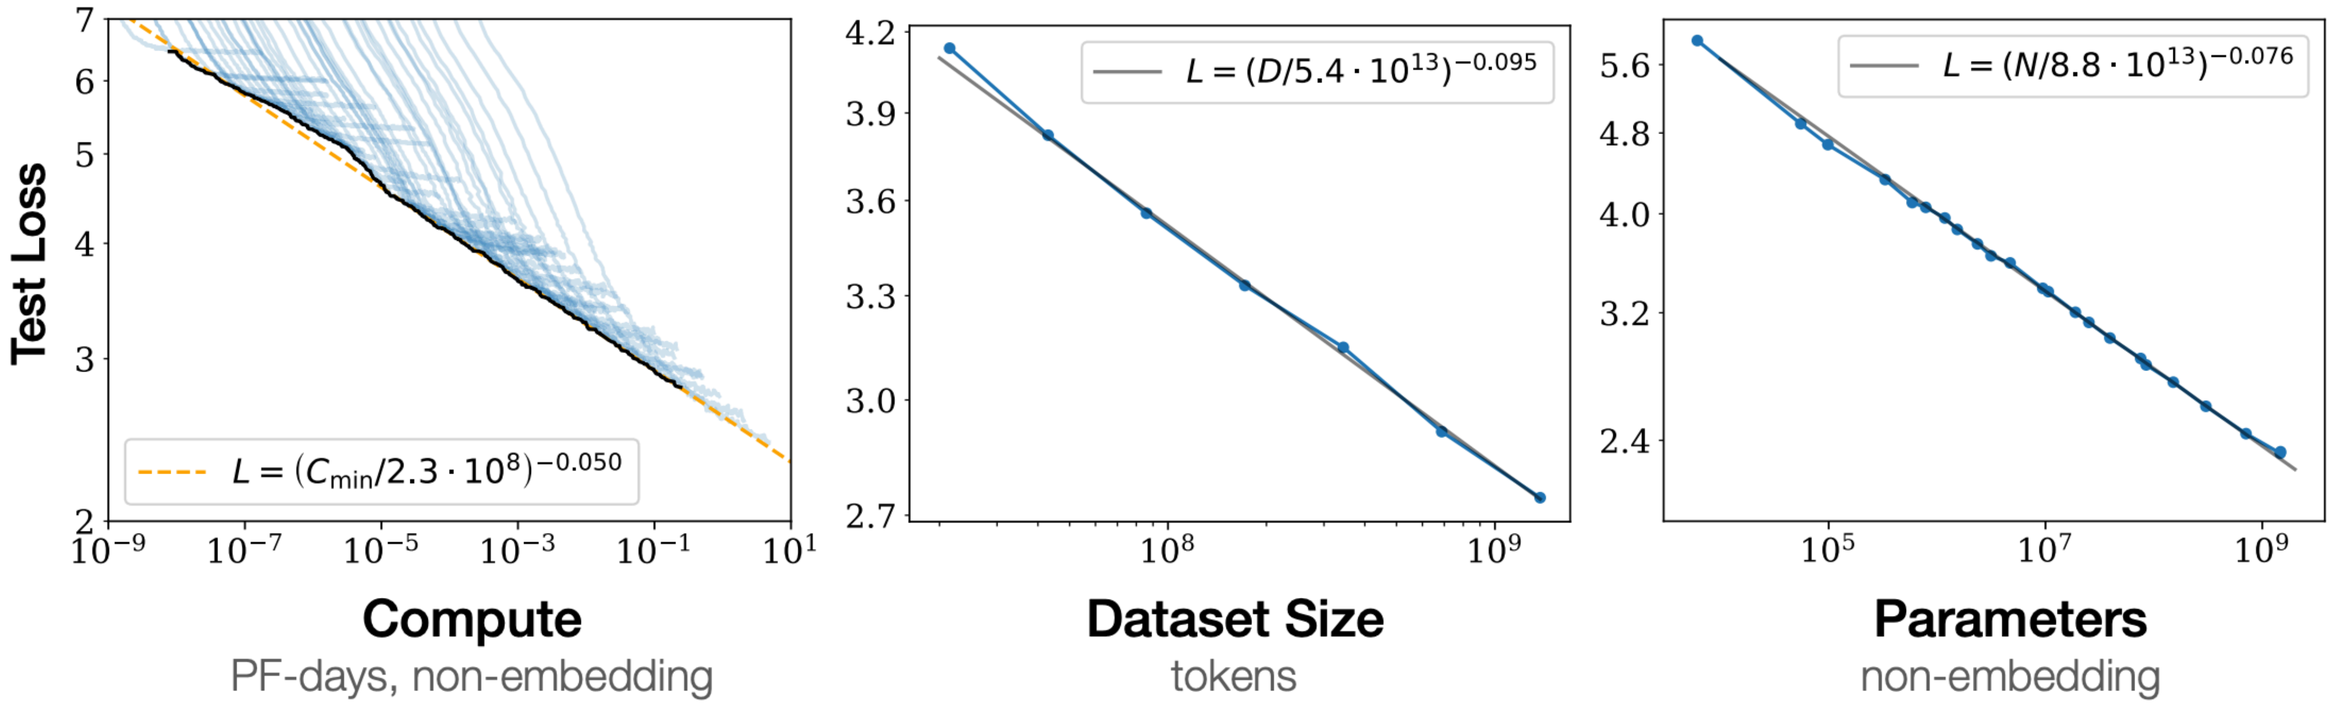

- Assuming no bottlenecks, expected test loss has power law relationship with each variable
- From smaller models, we can estimate how much compute, data, and model size is needed to achieve a particular test loss

Kaplan et al. (2020)

Model Scaling in the last two years

Model	Size (# Parameters)	Training Tokens
LaMDA (Thoppilan et al., 2022)	137 Billion	168 Billion
GPT-3 (Brown et al., 2020)	175 Billion	300 Billion
Jurassic (Lieber et al., 2021)	178 Billion	300 Billion
Gopher (Rae et al., 2021)	280 Billion	300 Billion
MT-NLG 530B (Smith et al., 2022)	530 Billion	270 Billion

What happens if we get these estimates wrong?

Oops!

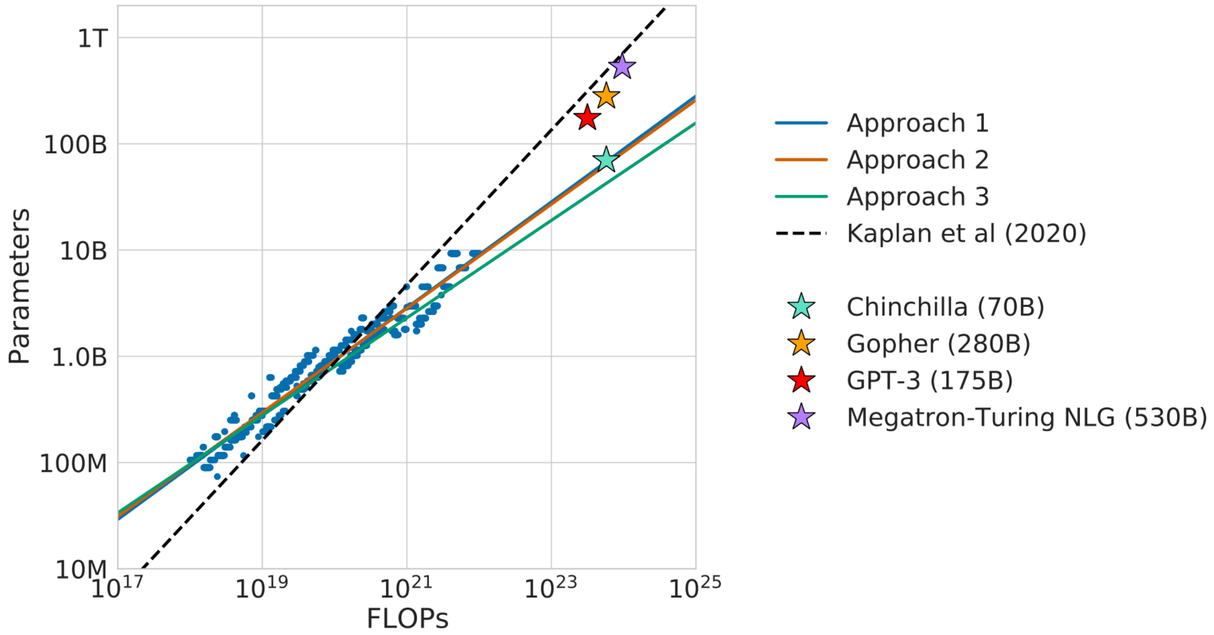

- Chinchilla authors founds that original works on model scaling had poorly estimated power laws
- New estimates showed that a 4x smaller model should be used for the compute budget
- Trained Gopher (280B) before finding this out!

Hoffmann et al. (2020)

Model Scaling in the last two years

Model	Size (# Parameters)	Training Tokens
LaMDA (Thoppilan et al., 2022)	137 Billion	168 Billion
GPT-3 (Brown et al., 2020)	175 Billion	300 Billion
Jurassic (Lieber et al., 2021)	178 Billion	300 Billion
Gopher (Rae et al., 2021)	280 Billion	300 Billion
MT-NLG 530B (Smith et al., 2022)	530 Billion	270 Billion
Chinchilla	70 Billion	1.4 Trillion

Chinchilla gets better performance than all of the above models on most common NLP benchmarks!

Smaller model, trained on much more data!

Hoffmann et al. (2020)

Recap

- Scale is necessary to achieve many of the emergent breakthroughs of the last few years
  - in-context learning, chain-of-thought reasoning, instruction learning
- Training at scale is very expensive
  - Potentially, months of training
  - millions of $$$$
- Scaling laws let us estimate the optimal model and dataset sizes for a fixed compute budget, so that we only have to do the training once!

References

- Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., ... & Amodei, D. (2020). Scaling laws for neural language models. arXiv preprint arXiv:2001.08361.

- Hoffmann, J., Borgeaud, S., Mensch, A., Buchatskaya, E., Cai, T., Rutherford, E., ... & Sifre, L. (2022). Training compute-optimal large language models. arXiv preprint arXiv:2203.15556.